

# Gradient Boosting Classification

```
from warnings import filterwarnings
filterwarnings('ignore')
```

## Reading the training dataset

```
import pandas as pd
df = pd.read_csv('train_titanic.csv')
df.head()
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

		Name	Sex	Age
SibSp	\			
0		Braund, Mr. Owen Harris	male	22.0
1				
1		Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0
1				
2		Heikkinen, Miss. Laina	female	26.0
0				
3		Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0
1				
4		Allen, Mr. William Henry	male	35.0
0				

	Parch		Ticket	Fare	Cabin	Embarked
0	0		A/5 21171	7.2500	NaN	S
1	0		PC 17599	71.2833	C85	C
2	0	STON/O2.	3101282	7.9250	NaN	S
3	0		113803	53.1000	C123	S
4	0		373450	8.0500	NaN	S

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     891 non-null   int64
1   Survived        891 non-null   int64
2   Pclass          891 non-null   int64
```

```
3   Name      891 non-null  object
4   Sex       891 non-null  object
5   Age       714 non-null  float64
6   SibSp     891 non-null  int64
7   Parch     891 non-null  int64
8   Ticket    891 non-null  object
9   Fare      891 non-null  float64
10  Cabin     204 non-null  object
11  Embarked  889 non-null  object
```

dtypes: float64(2), int64(5), object(5)

memory usage: 83.7+ KB

```
df.nunique()
```

```
PassengerId    891
Survived        2
Pclass         3
Name           891
Sex            2
Age            88
SibSp          7
Parch          7
Ticket        681
Fare          248
Cabin         147
Embarked       3
dtype: int64
```

```
s = df.isna().sum()
s
```

```
PassengerId    0
Survived       0
Pclass         0
Name           0
Sex            0
Age           177
SibSp          0
Parch          0
Ticket         0
Fare           0
Cabin         687
Embarked       2
dtype: int64
```

```
s[s>0]
```

```
Age           177
Cabin         687
Embarked       2
dtype: int64
```

## Seprating X and Y

```
X = df.drop(labels=['PassengerId', 'Name', 'Ticket', 'Survived'], axis=1)
Y = df[['Survived']]
```

```
X.head()
```

	Pclass	Sex	Age	SibSp	Parch	Fare	Cabin	Embarked
0	3	male	22.0	1	0	7.2500	NaN	S
1	1	female	38.0	1	0	71.2833	C85	C
2	3	female	26.0	0	0	7.9250	NaN	S
3	1	female	35.0	1	0	53.1000	C123	S
4	3	male	35.0	0	0	8.0500	NaN	S

```
Y.head()
```

	Survived
0	0
1	1
2	1
3	1
4	0

## Creating a pipline for X preprocessing

```
cat = list(X.columns[X.dtypes=='object'])
cat

['Sex', 'Cabin', 'Embarked']

con = list(X.columns[X.dtypes!='object'])
con

['Pclass', 'Age', 'SibSp', 'Parch', 'Fare']

cat1 = ['Cabin']
cat2 = ['Sex', 'Embarked']
```

## Preprocessing Pipeline

```
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.compose import ColumnTransformer

# Num Pipeline
num_pipe =
Pipeline(steps=[('impute', SimpleImputer(strategy='median')),
                ('scaler', StandardScaler())])

# Pipeline for cabin column
cat_pipe1 =
```

```
Pipeline(steps=[('imputer',SimpleImputer(strategy='constant',fill_value='Unknown'))],
```

```
('Ohe',OneHotEncoder(handle_unknown='ignore'))])
```

```
# Pipeline for remaining features
```

```
cat_pipe2 =
```

```
Pipeline(steps=[('imuter',SimpleImputer(strategy='most_frequent'))],
```

```
('Ohe',OneHotEncoder(handle_unknown='ignore'))])
```

```
# Compose
```

```
from sklearn.compose import ColumnTransformer
```

```
pre = ColumnTransformer([('num',num_pipe,con),  
                          ('cat1',cat_pipe1,cat1),  
                          ('cat2',cat_pipe2,cat2)])
```

```
X_pre = pre.fit_transform(X).toarray()
```

```
X_pre
```

```
array([[ 0.82737724, -0.56573646,  0.43279337, ...,  0.        ,  
        0.        ,  1.        ],  
       [-1.56610693,  0.66386103,  0.43279337, ...,  1.        ,  
        0.        ,  0.        ],  
       [ 0.82737724, -0.25833709, -0.4745452 , ...,  0.        ,  
        0.        ,  1.        ],  
       ...,  
       [ 0.82737724, -0.1046374 ,  0.43279337, ...,  0.        ,  
        0.        ,  1.        ],  
       [-1.56610693, -0.25833709, -0.4745452 , ...,  1.        ,  
        0.        ,  0.        ],  
       [ 0.82737724,  0.20276197, -0.4745452 , ...,  0.        ,  
        1.        ,  0.        ]])
```

```
cols = pre.get_feature_names_out()
```

```
cols
```

```
array(['num__Pclass', 'num__Age', 'num__SibSp', 'num__Parch',  
      'num__Fare',  
      'cat1__Cabin_A10', 'cat1__Cabin_A14', 'cat1__Cabin_A16',  
      'cat1__Cabin_A19', 'cat1__Cabin_A20', 'cat1__Cabin_A23',  
      'cat1__Cabin_A24', 'cat1__Cabin_A26', 'cat1__Cabin_A31',  
      'cat1__Cabin_A32', 'cat1__Cabin_A34', 'cat1__Cabin_A36',  
      'cat1__Cabin_A5', 'cat1__Cabin_A6', 'cat1__Cabin_A7',  
      'cat1__Cabin_B101', 'cat1__Cabin_B102', 'cat1__Cabin_B18',  
      'cat1__Cabin_B19', 'cat1__Cabin_B20', 'cat1__Cabin_B22',  
      'cat1__Cabin_B28', 'cat1__Cabin_B3', 'cat1__Cabin_B30',  
      'cat1__Cabin_B35', 'cat1__Cabin_B37', 'cat1__Cabin_B38',  
      'cat1__Cabin_B39', 'cat1__Cabin_B4', 'cat1__Cabin_B41',  
      'cat1__Cabin_B42', 'cat1__Cabin_B49', 'cat1__Cabin_B5',  
      'cat1__Cabin_B50', 'cat1__Cabin_B51 B53 B55',
```

```

'cat1__Cabin_B57 B59 B63 B66', 'cat1__Cabin_B58 B60',
'cat1__Cabin_B69', 'cat1__Cabin_B71', 'cat1__Cabin_B73',
'cat1__Cabin_B77', 'cat1__Cabin_B78', 'cat1__Cabin_B79',
'cat1__Cabin_B80', 'cat1__Cabin_B82 B84', 'cat1__Cabin_B86',
'cat1__Cabin_B94', 'cat1__Cabin_B96 B98', 'cat1__Cabin_C101',
'cat1__Cabin_C103', 'cat1__Cabin_C104', 'cat1__Cabin_C106',
'cat1__Cabin_C110', 'cat1__Cabin_C111', 'cat1__Cabin_C118',
'cat1__Cabin_C123', 'cat1__Cabin_C124', 'cat1__Cabin_C125',
'cat1__Cabin_C126', 'cat1__Cabin_C128', 'cat1__Cabin_C148',
'cat1__Cabin_C2', 'cat1__Cabin_C22 C26', 'cat1__Cabin_C23 C25
C27',
'cat1__Cabin_C30', 'cat1__Cabin_C32', 'cat1__Cabin_C45',
'cat1__Cabin_C46', 'cat1__Cabin_C47', 'cat1__Cabin_C49',
'cat1__Cabin_C50', 'cat1__Cabin_C52', 'cat1__Cabin_C54',
'cat1__Cabin_C62 C64', 'cat1__Cabin_C65', 'cat1__Cabin_C68',
'cat1__Cabin_C7', 'cat1__Cabin_C70', 'cat1__Cabin_C78',
'cat1__Cabin_C82', 'cat1__Cabin_C83', 'cat1__Cabin_C85',
'cat1__Cabin_C86', 'cat1__Cabin_C87', 'cat1__Cabin_C90',
'cat1__Cabin_C91', 'cat1__Cabin_C92', 'cat1__Cabin_C93',
'cat1__Cabin_C95', 'cat1__Cabin_C99', 'cat1__Cabin_D',
'cat1__Cabin_D10 D12', 'cat1__Cabin_D11', 'cat1__Cabin_D15',
'cat1__Cabin_D17', 'cat1__Cabin_D19', 'cat1__Cabin_D20',
'cat1__Cabin_D21', 'cat1__Cabin_D26', 'cat1__Cabin_D28',
'cat1__Cabin_D30', 'cat1__Cabin_D33', 'cat1__Cabin_D35',
'cat1__Cabin_D36', 'cat1__Cabin_D37', 'cat1__Cabin_D45',
'cat1__Cabin_D46', 'cat1__Cabin_D47', 'cat1__Cabin_D48',
'cat1__Cabin_D49', 'cat1__Cabin_D50', 'cat1__Cabin_D56',
'cat1__Cabin_D6', 'cat1__Cabin_D7', 'cat1__Cabin_D9',
'cat1__Cabin_E10', 'cat1__Cabin_E101', 'cat1__Cabin_E12',
'cat1__Cabin_E121', 'cat1__Cabin_E17', 'cat1__Cabin_E24',
'cat1__Cabin_E25', 'cat1__Cabin_E31', 'cat1__Cabin_E33',
'cat1__Cabin_E34', 'cat1__Cabin_E36', 'cat1__Cabin_E38',
'cat1__Cabin_E40', 'cat1__Cabin_E44', 'cat1__Cabin_E46',
'cat1__Cabin_E49', 'cat1__Cabin_E50', 'cat1__Cabin_E58',
'cat1__Cabin_E63', 'cat1__Cabin_E67', 'cat1__Cabin_E68',
'cat1__Cabin_E77', 'cat1__Cabin_E8', 'cat1__Cabin_F E69',
'cat1__Cabin_F G63', 'cat1__Cabin_F G73', 'cat1__Cabin_F2',
'cat1__Cabin_F33', 'cat1__Cabin_F38', 'cat1__Cabin_F4',
'cat1__Cabin_G6', 'cat1__Cabin_T', 'cat1__Cabin_Unknown',
'cat2__Sex_female', 'cat2__Sex_male', 'cat2__Embarked_C',
'cat2__Embarked_Q', 'cat2__Embarked_S'], dtype=object)

```

```
len(cols)
```

```
158
```

```

X_pre = pd.DataFrame(X_pre, columns=cols)
X_pre.head()

```

	num__Pclass	num__Age	num__SibSp	num__Parch	num__Fare
cat1__Cabin_A10 \					
0	0.827377	-0.565736	0.432793	-0.473674	-0.502445
0.0					
1	-1.566107	0.663861	0.432793	-0.473674	0.786845
0.0					
2	0.827377	-0.258337	-0.474545	-0.473674	-0.488854
0.0					
3	-1.566107	0.433312	0.432793	-0.473674	0.420730
0.0					
4	0.827377	0.433312	-0.474545	-0.473674	-0.486337
0.0					

	cat1__Cabin_A14	cat1__Cabin_A16	cat1__Cabin_A19	cat1__Cabin_A20
... \				
0	0.0	0.0	0.0	0.0
...				
1	0.0	0.0	0.0	0.0
...				
2	0.0	0.0	0.0	0.0
...				
3	0.0	0.0	0.0	0.0
...				
4	0.0	0.0	0.0	0.0
...				

	cat1__Cabin_F38	cat1__Cabin_F4	cat1__Cabin_G6	cat1__Cabin_T	\
0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	

	cat1__Cabin_Unknown	cat2__Sex_female	cat2__Sex_male
cat2__Embarked_C \			
0	1.0	0.0	1.0
0.0			
1	0.0	1.0	0.0
1.0			
2	1.0	1.0	0.0
0.0			
3	0.0	1.0	0.0
0.0			
4	1.0	0.0	1.0
0.0			

	cat2__Embarked_Q	cat2__Embarked_S
0	0.0	1.0
1	0.0	0.0
2	0.0	1.0

```
3          0.0          1.0
4          0.0          1.0
```

```
[5 rows x 158 columns]
```

```
X_pre.isna().sum()
```

```
num__Pclass      0
num__Age          0
num__SibSp        0
num__Parch        0
num__Fare         0
..
cat2__Sex_female  0
cat2__Sex_male    0
cat2__Embarked_C  0
cat2__Embarked_Q  0
cat2__Embarked_S  0
Length: 158, dtype: int64
```

```
Y.value_counts()
```

```
Survived
0          549
1          342
Name: count, dtype: int64
```

## Train test split

```
from sklearn.model_selection import train_test_split
xtrain, xtest, ytrain, ytest = train_test_split(X_pre, Y,
test_size=0.2, random_state=63)

xtrain.shape
(712, 158)

xtest.shape
(179, 158)
```

## Model creation

```
from sklearn.ensemble import GradientBoostingClassifier
model = GradientBoostingClassifier(random_state=42)
model.fit(xtrain, ytrain)

GradientBoostingClassifier(random_state=42)

# Training Accuracy
model.score(xtrain, ytrain)
```

```
0.8960674157303371
```

```
# Testing
```

```
model.score(xtest, ytest)
```

```
0.8268156424581006
```

## Hyperparameter tuning

```
params = {'learning_rate':[0.001, 0.01, 0.05, 0.1],  
          'n_estimators':[10,50,100,200],  
          'max_depth':[3,4,5,6,7,8,9,10],  
          'min_samples_split':[6,7,8,9,10]}
```

```
from sklearn.model_selection import RandomizedSearchCV  
gbc = GradientBoostingClassifier(random_state=42)  
rscv = RandomizedSearchCV(gbc, param_distributions=params, cv=5,  
scoring='f1')  
rscv.fit(xtrain,ytrain)
```

```
RandomizedSearchCV(cv=5,  
estimator=GradientBoostingClassifier(random_state=42),  
                  param_distributions={'learning_rate': [0.001, 0.01,  
0.05,  
                                0.1],  
                                'max_depth': [3, 4, 5, 6, 7,  
8, 9, 10],  
                                'min_samples_split': [6, 7, 8,  
9, 10],  
                                'n_estimators': [10, 50, 100,  
200]}},  
                  scoring='f1')
```

```
rscv.best_params_
```

```
{'n_estimators': 200,  
 'min_samples_split': 9,  
 'max_depth': 5,  
 'learning_rate': 0.1}
```

```
rscv.best_score_
```

```
0.7394691774362689
```

```
best_gbc = rscv.best_estimator_  
best_gbc
```

```
GradientBoostingClassifier(max_depth=5, min_samples_split=9,  
n_estimators=200,  
                           random_state=42)
```



## Evaluate best gbc

```
best_gbc.score(xtrain,ytrain)
0.9325842696629213
best_gbc.score(xtest, ytest)
0.8324022346368715
```

## Predicting training and testing data

```
ypred_tr = best_gbc.predict(xtrain)
ypred_ts = best_gbc.predict(xtest)

ypred_tr[0:5]
array([0, 0, 0, 0, 0], dtype=int64)

ytrain.head()
  Survived
610      0
728      0
90       0
509      1
834      0


ypred_ts[0:5]
array([0, 0, 0, 1, 0], dtype=int64)

ytest.head()
  Survived
789      0
823      1
4       0
3       1
721      0
```

## Confusion Matrix

```
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
cf = confusion_matrix(ytest, ypred_ts)
cfd = ConfusionMatrixDisplay(cf,display_labels=best_gbc.classes_)
cfd.plot()

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x20735f3c410>
```



## Perform classification report

```
from sklearn.metrics import classification_report
print(classification_report(ytest, ypred_ts))
```

	precision	recall	f1-score	support
0	0.81	0.93	0.86	103
1	0.88	0.70	0.78	76
accuracy			0.83	179
macro avg	0.85	0.81	0.82	179
weighted avg	0.84	0.83	0.83	179

## Use the model for predictions

```
xnew = pd.read_csv('test_titanic.csv')
xnew.head()
```

	PassengerId	Pclass	Name
Sex \			
0	892	3	Kelly, Mr. James
male			
1	893	3	Wilkes, Mrs. James (Ellen Needs)

```
female
2      894      2      Myles, Mr. Thomas Francis
male
3      895      3      Wirz, Mr. Albert
male
4      896      3  Hirvonen, Mrs. Alexander (Helga E Lindqvist)
female
```

	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	34.5	0	0	330911	7.8292	NaN	Q
1	47.0	1	0	363272	7.0000	NaN	S
2	62.0	0	0	240276	9.6875	NaN	Q
3	27.0	0	0	315154	8.6625	NaN	S
4	22.0	1	1	3101298	12.2875	NaN	S

```
xnew.isna().sum()
```

```
PassengerId      0
Pclass            0
Name              0
Sex               0
Age              86
SibSp             0
Parch            0
Ticket           0
Fare              1
Cabin           327
Embarked         0
dtype: int64
```

```
xnew_pre = pre.transform(xnew).toarray()
xnew_pre
```

```
array([[ 0.82737724,  0.39488658, -0.4745452 , ...,  0.        ,
         1.        ,  0.        ],
       [ 0.82737724,  1.35550962,  0.43279337, ...,  0.        ,
         0.        ,  1.        ],
       [-0.36936484,  2.50825727, -0.4745452 , ...,  0.        ,
         1.        ,  0.        ],
       ...,
       [ 0.82737724,  0.70228595, -0.4745452 , ...,  0.        ,
         0.        ,  1.        ],
       [ 0.82737724, -0.1046374 , -0.4745452 , ...,  0.        ,
         0.        ,  1.        ],
       [ 0.82737724, -0.1046374 ,  0.43279337, ...,  1.        ,
         0.        ,  0.        ]])
```

```
xnew_pre = pd.DataFrame(xnew_pre,columns=cols)
xnew_pre
```

	num_Pclass	num_Age	num_SibSp	num_Parch	num_Fare	\
0	0.827377	0.394887	-0.474545	-0.473674	-0.490783	
1	0.827377	1.355510	0.432793	-0.473674	-0.507479	
2	-0.369365	2.508257	-0.474545	-0.473674	-0.453367	
3	0.827377	-0.181487	-0.474545	-0.473674	-0.474005	
4	0.827377	-0.565736	0.432793	0.767630	-0.401017	
..	...	...	...	...	...	
413	0.827377	-0.104637	-0.474545	-0.473674	-0.486337	
414	-1.566107	0.740711	-0.474545	-0.473674	1.544246	
415	0.827377	0.702286	-0.474545	-0.473674	-0.502445	
416	0.827377	-0.104637	-0.474545	-0.473674	-0.486337	
417	0.827377	-0.104637	0.432793	0.767630	-0.198244	
	cat1_Cabin_A10	cat1_Cabin_A14	cat1_Cabin_A16	cat1_Cabin_A19	\	
0	0.0	0.0	0.0	0.0		
0.0						
1	0.0	0.0	0.0	0.0		
0.0						
2	0.0	0.0	0.0	0.0		
0.0						
3	0.0	0.0	0.0	0.0		
0.0						
4	0.0	0.0	0.0	0.0		
0.0						
..	...	...	...	...		..
.						
413	0.0	0.0	0.0	0.0		
0.0						
414	0.0	0.0	0.0	0.0		
0.0						
415	0.0	0.0	0.0	0.0		
0.0						
416	0.0	0.0	0.0	0.0		
0.0						
417	0.0	0.0	0.0	0.0		
0.0						
	cat1_Cabin_A20	...	cat1_Cabin_F38	cat1_Cabin_F4	cat1_Cabin_G6	\
0	0.0	...	0.0	0.0	0.0	
0.0						
1	0.0	...	0.0	0.0	0.0	
0.0						
2	0.0	...	0.0	0.0	0.0	
0.0						
3	0.0	...	0.0	0.0	0.0	
0.0						
4	0.0	...	0.0	0.0	0.0	
0.0						

...	...	...	...
...			
413	0.0	...	0.0
0.0			
414	0.0	...	0.0
0.0			
415	0.0	...	0.0
0.0			
416	0.0	...	0.0
0.0			
417	0.0	...	0.0
0.0			

	cat1__Cabin_T	cat1__Cabin_Unknown	cat2__Sex_female
cat2__Sex_male \			
0	0.0	1.0	0.0
1.0			
1	0.0	1.0	1.0
0.0			
2	0.0	1.0	0.0
1.0			
3	0.0	1.0	0.0
1.0			
4	0.0	1.0	1.0
0.0			
...	...	...	...
...			
413	0.0	1.0	0.0
1.0			
414	0.0	0.0	1.0
0.0			
415	0.0	1.0	0.0
1.0			
416	0.0	1.0	0.0
1.0			
417	0.0	1.0	0.0
1.0			

	cat2__Embarked_C	cat2__Embarked_Q	cat2__Embarked_S
0	0.0	1.0	0.0
1	0.0	0.0	1.0
2	0.0	1.0	0.0
3	0.0	0.0	1.0
4	0.0	0.0	1.0
...	...	...	...
413	0.0	0.0	1.0
414	1.0	0.0	0.0
415	0.0	0.0	1.0
416	0.0	0.0	1.0
417	1.0	0.0	0.0

```
[418 rows x 158 columns]
```

## Use the model for predictions

```
pred = best_gbc.predict(xnew_pre)
pred
```

```
array([0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 0,
0,
      1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0,
1,
      1, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1,
1,
      1, 0, 1, 1, 1, 0, 0, 1, 1, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1,
0,
      1, 1, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,
1,
      0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0,
1,
      1, 0, 1, 1, 0, 1, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1,
0,
      0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0,
0,
      1, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 1,
1,
      0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0,
1,
      0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0,
1,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1,
1,
      0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1,
0,
      1, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1,
1,
      0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1,
0,
      1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0,
1,
      0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 1, 1, 1, 1, 0, 0, 1, 0,
0],
      dtype=int64)
```

```
len(pred)
```

## Saving Predictions in Dataframe

```
df_pred = xnew[['PassengerId']]
df_pred
```

	PassengerId
0	892
1	893
2	894
3	895
4	896
...	...
413	1305
414	1306
415	1307
416	1308
417	1309

```
[418 rows x 1 columns]
```

```
df_pred['Survived_Pred'] = pred
df_pred
```

	PassengerId	Survived_Pred
0	892	0
1	893	0
2	894	0
3	895	0
4	896	0
...	...	...
413	1305	0
414	1306	1
415	1307	0
416	1308	0
417	1309	0

```
[418 rows x 2 columns]
```

```
df_pred['Survived_Pred'].value_counts()
```

```
Survived_Pred
```


```
0    283
```

```
1    135
```

```
Name: count, dtype: int64
```

```
df_pred['Survived_Pred'].value_counts().plot(kind='bar')
```

```
<Axes: xlabel='Survived_Pred'>
```



## Save Predictions to CSV

```
df_pred.to_csv('Survived_Predictions.csv', index = False)
```